

Save it as a skill

You have spent an afternoon working out how your executive writes, travels and decides. Here is how to stop that document going stale in a folder, and turn it into something that applies itself.

01 Executive Playbook

02 Executive Voice

03 LinkedIn

What this is

Three of the SortedEA skills end by handing you a document about your principal. This guide turns each of those documents into a skill of its own, so it works without being remembered.

The pattern is the same in all three cases. You run a skill, it interviews you and reads what you give it, and you end up with a filled-in file: a playbook, a voice profile, a set of LinkedIn house rules. That file is genuinely useful. It is also a file, which means somebody has to remember it exists and open it at the right moment.

A skill does not have that problem. Once installed, it reads itself in whenever the work matches, and applies what you wrote without anyone thinking about it. The document becomes infrastructure.

This takes about 20 minutes per skill and needs no technical background. If you have ever renamed a folder and made a zip file, you can do all of it.

- 1 What this is
- 2 How a skill is put together the shared mechanics
- 3 Executive Playbook → a skill
- 4 Executive Voice → a skill
- 5 LinkedIn → a skill
- 6 The five privacy rules read before packaging
- 7 When it does not work
- 8 The whole thing on one page

BEFORE YOU START

Skills work on every Claude plan, but they need code execution switched on. On Free, Pro and Max that is Settings, then Capabilities. On Team and Enterprise it is Organization settings, then Skills.

If the option to create or upload a skill is missing entirely, your organisation has probably turned off user-created skills. That is an administrator setting, not a fault. Ask before you spend an afternoon on it.

TWO

How a skill is put together

This section applies to all three. Read it once, then go to the one you are building.

The anatomy

A skill is a folder. Inside it sits a file called `SKILL.md`, and optionally a `references` folder holding material that is only needed sometimes.

```
voice-jchen/
├── SKILL.md           the rules, read every time
└── references/
    ├── samples.md     read only when needed
    └── log.md          read only when needed
```

That split is the single most useful thing to understand. Everything in `SKILL.md` loads into the conversation every time the skill fires. Reference files sit on disk and are opened only when the work calls for them.

So `SKILL.md` holds the rules you want applied to every piece of work. Long lists, verbatim examples and running logs go into references, where they cost nothing until they are wanted.

The two limits, and the trap

The top of `SKILL.md` carries a short block of settings. It looks like this, three dashes above and below:

```
---
name: voice-jchen
description: How J Chen writes. Apply when drafting any email or note
            going out under their name, or when checking a draft sounds like them.
---
```

FIELD	LIMIT	WHAT CATCHES PEOPLE OUT
name	64 characters	Lowercase letters, numbers and hyphens only. No spaces, no capitals, no underscores. It cannot contain the words <i>claude</i> or <i>anthropic</i> , which are reserved.
description	200 characters	The only thing Claude reads when deciding whether to use the skill at all. It has to say what the skill is <i>and</i> when to use it.

FOLDER NAME AND SKILL NAME

You name the same thing twice, in the folder and in the `name` field. The field is the one that counts, and nothing enforces a match. Match them anyway, or in six months you will not know which file to open when you want to change something.

THE COLON TRAP

Do not put a colon in the description. An unquoted colon breaks the settings block, and the skill fails to load with no error message and no warning. It simply never fires, and you spend twenty minutes wondering why.

If you need one, wrap the whole description in quotation marks. Easier: rewrite the sentence without it.

ON THE 200

You may see 1,024 quoted elsewhere. That figure is from Anthropic's developer documentation and applies to the developer platform. The Claude app enforces 200. Write to 200 and it works everywhere; anything that will not fit belongs in the body, where there is no limit.

Write the description for what you actually say

The description is a matching problem, not a summary. Claude compares what you typed against it, so it needs the words that come out of your mouth on a normal Tuesday.

If you say "*draft a reply to Marcus*", then "drafting", "reply" and "in their name" need to be in there. If you say "*what does the playbook say about her travel*", then "playbook" and "travel" do. Nobody types "apply the voice profile", so do not build the description around that phrase.

Naming, when there is more than one principal

Use a prefix that survives a second person: `voice-jchen`, `playbook-jchen`, `linkedin-jchen`. Then `voice-mrivera` when you pick up a second principal.

Do not call it `voice` or `my-boss`. The day you support two people, neither you nor Claude will know which is which, and the descriptions will collide at the worst possible moment.

If the file might be seen by someone who should not connect a document to a person, use initials rather than a full name in the folder.

Turning a document into instructions

Same content, imperative mood. The document describes. The skill instructs.

THE DOCUMENT SAYS

Mean sentence length 14 words in internal email, measured across 31 samples.

THE SKILL SAYS

Keep internal email sentences short, around 14 words. Long clause-heavy sentences do not sound like them.

Keep the numbers. "Prefers short sentences" is not actionable. "Around 14 words, from 31 emails" is. The measurements are what make this better than a hunch, so do not smooth them out into adjectives.

Keep the strength labels. A skill that applies every tendency at full force produces a parody. Carry the labels across, and explain them at the top of the body:

HARD never violate. **STRONG** most of the time. Everything else is a light preference that context can override.

Point at the reference files. They are only read if something tells Claude to read them. One line does it:

When a draft feels wrong, read `references/samples.md` and match the rhythm of the real examples rather than re-reading these rules.

Two lines to end every one of them with

Matching the voice or the preference never justifies changing a fact, softening a decline, or committing them to something they did not say.

Do not force every pattern into every piece of work. If it reads as an imitation rather than as them, use fewer of the patterns, not more.

Packaging and installing

- 1 **Zip the folder, not the files.** The zip must contain `voice-jchen/` with `SKILL.md` inside it. A zip with a loose `SKILL.md` at the top level will be rejected.
- 2 **In the Claude app:** Customize, then Skills, then Add, then "Upload a skill". Choose the zip.
- 3 **Check it is switched on** in your skills list afterwards. Uploading and enabling are two different things.
- 4 **In Claude Code or a similar tool:** drop the folder into `~/.claude/skills/` for personal use, or `.claude/skills/` inside a project.

DO NOT PASTE IT IN

The Add menu also offers "Write skill instructions", where you paste text straight in. That works for a single-file skill, but it flattens everything into one block and you lose the reference files. Upload the zip instead.

Test it before you rely on it

Give it something you have already done by hand and compare the two closely, looking for small differences rather than obvious failures. Then give it an awkward case: a missing detail, an unusual request, something the document does not cover. You need to know how it behaves at the edges before it meets one while you are not watching.

Keeping it current

All three have the same maintenance loop, and it is the reason a skill gets better rather than staler.

Log the corrections. When your principal changes something you drafted, or tells you a preference has moved, write down what you had, what they changed it to, and the rule behind it. Put it in `references/log.md`.

Once a quarter, promote. Anything in the log that has happened three times becomes a rule in `SKILL.md`. Fifteen minutes.

Executive Playbook → a skill

The operating file becomes a skill. The private annex never does.

What you have after the skill runs

playbook.md + playbook-private.md

Two files, deliberately separate. The operating file covers how they work: hours, meeting style, who gets straight through, seat and hotel preferences, how they want bad news delivered. The private annex holds identity and personal detail: date of birth, home address, family, medical contacts, loyalty numbers.

Only the operating file becomes a skill. That separation is the entire reason the playbook was built in two parts.

Name it

```
name: playbook-jchen
description: How J Chen works. Read before scheduling, booking travel,
  drafting on their behalf or answering what they prefer.
```

What goes in SKILL.md

SECTION	WHY IT BELONGS IN THE ALWAYS-LOADED PART
Scheduling rules	Meeting length, dead zones, non-negotiables, buffers, and who gets through against who is deflected. The most-used section by a distance.
Daily routine	Working hours, protected thinking time, when they are reachable, weekend rules.
Communication	Channel by urgency, what you may send without checking, what always comes back to them first.
Travel preferences	Airline, cabin, seat, hotels by city, departure habits, buffers. Preferences only. Numbers stay in the annex.
Work profile	Meeting style, how they decide, what they consider a waste of time, what they always want to be asked about first.
Emergency procedure	Who you call first, who is told in what order, what you may cancel without asking. The procedure, not the medical contacts.
Confidentiality rules	Including what may and may not go into AI tools. Worth having in front of Claude every time.

What goes in references

`references/preferences.md` for the long tail: restaurants by occasion, gifts that landed well, favourite vendors, food and drink detail, the small things that make a bad day better. Genuinely useful, rarely needed mid-task.

`references/log.md` for corrections.

NEVER INCLUDE

Nothing from the private annex goes into the skill. Not the date of birth, not the home address, not the family names, not the doctor, not the loyalty numbers, not the passport details.

The skill may say *"passport and payment details are in the annex, ask me"*. It may not contain them. If a task genuinely needs one, you supply it by hand at that moment.

Passwords do not go in either file, or any file. That has not changed.

Carry the tags across

The playbook marks each line `[said]`, `[observed]` or `[inherited]`. Keep them, and tell the skill what to do with them:

Act on `[said]` directly. For `[observed]`, act but flag it when the stakes are high: "the playbook has this as observed rather than confirmed, worth a quick check?" Treat `[inherited]` as a hypothesis.

And carry across the instruction that matters most: **when the playbook does not cover something, say so and ask, rather than inventing a plausible preference**. A playbook that quietly fills its own gaps becomes fiction that everybody trusts.

The test

Ask it five things you already know the answer to. Then ask it something the playbook does not cover, and check that it says so instead of guessing.

Executive Voice → a skill

This is the one that pays back fastest, because drafting happens every day.

What you have after the skill runs

`voice-profile.md`

Measurements taken from 25 to 40 real pieces of their writing, plus the answers to about 20 questions covering what writing cannot show. Sentence lengths by register, greeting and sign-off counts, signature vocabulary, how they decline, how they chase, what never appears anywhere in the corpus.

Name it

```
name: voice-jchen
description: How J Chen writes. Apply when drafting any email, message
  or note going out under their name, or checking a draft sounds like them.
```

What goes in SKILL.md

- The three-line summary
- The measurements, with their numbers intact
- Openings and closings, with the counts
- Signature vocabulary, and the words they avoid
- Mechanics: punctuation they use, punctuation that never appears, bullets or prose, where the ask sits
- How they handle declines, chases, disagreement, bad news, praise
- The register table: what changes between a direct report and the board, internal and external, good news and bad
- What never appears in the corpus, which is the strongest evidence you have
- The hard nos

What goes in references

`references/samples.md` holds every verbatim example, stripped of names and figures. This is the most valuable file in the whole set and the one most people leave out.

When a draft comes out wrong, reading three real sentences fixes it faster than re-reading any amount of analysis. But those samples do not need to be in context for every message, which is exactly what references are for.

`references/log.md` for the corrections.

NEVER INCLUDE

Strip the samples again at packaging time. You did it once when building the profile. Do it again now, because samples are the part most likely to carry a client name, a figure, or a sentence about a third person.

A voice sample needs the shape, not the content. "*Can you get me the [figure] before [day]?*" carries everything useful and discloses nothing.

The test that actually works

Write three short pieces with the skill on. Mix them with three real ones. Show a colleague who knows your principal well and ask which are real.

If they pick yours out immediately, the profile is not finished, and their answer to "*what gave it away?*" is the rule you are missing.

THE BETTER, SLOWER TEST

Send them a draft and watch what they change. If they only change facts and content, the voice is working. If they rewrite phrasing, every rewrite is a rule you missed. Log it.

Keep doing this after the skill is finished. Corrections are the highest quality evidence there is and they arrive free, forever.

LinkedIn → a skill

Not how LinkedIn works. What *this* principal will and will not say in public.

What you have after the skill runs

`linkedin-house-rules.md`

The decisions you and your principal have made about their presence: what they post about, what they will never touch, who publishes, what the approval route is, and the shape their posts take.

The general skill teaches the platform and the policy. This one records the specific arrangement you have with this specific person.

Name it

```
name: linkedin-jchen
description: J Chen's LinkedIn house rules. Use when drafting a post
            or comment for them, or deciding whether something should be posted.
```

What goes in SKILL.md

SECTION	CONTENTS
Territory	The three or four subjects they have real authority on and are willing to talk about publicly.
The no-go list	Named subjects they will not comment on. Usually the most specific and most useful section in the file.
Post voice	How their public writing differs from their email. It usually does. Reference the voice skill for the underlying register and record only the delta.
Post shapes	The two or three formats that work for them. An observation from a meeting, a position taken, a piece of news with a view attached.
Publication route	Who publishes and how. State plainly that you do not have their login and are not going to.
Comment policy	Whose posts they engage with, and whether comments are drafted or theirs alone.
The check	Is it theirs to say, does it commit the organisation, does it name someone who has not agreed, would they have to delete it.

What goes in references

`references/posts.md` for their posts that worked, and the ones that did not, with a line on why.
`references/log.md` for what they changed.

PUT THE SUBSTANCE RULE IN THE BODY

Since 2026 LinkedIn names "AI slop" as a category it demotes, describing it as content that "may sound polished on the surface but lacks a clear point of view, unique perspective, or substance". Readers can report it, and enough reports surface a warning in the author's own analytics.

So the failure mode has inverted. It is no longer sounding wrong. It is **sounding fine and saying nothing**, which is both the easiest thing to produce and the thing the platform now suppresses.

Put this instruction in the skill, in these words or your own: *if there is no specific claim, number, named example or genuine opinion in this, do not write the post. Say there is nothing to post this week.*

NEVER INCLUDE

Their login. Ever, anywhere, in any file.

Signing in as your principal breaches LinkedIn's User Agreement, which states plainly that you will not share your account, and separately prohibits using another person's account. The account holder is responsible for everything that happens through the account, so the exposure lands on them as much as on you.

Record the publication route instead: they publish, or an authorised scheduling tool does, using LinkedIn's official permission flow that they approve themselves and can revoke.

The test

Draft three posts. Expect at least one of them to fail your own no-go check, and if none of them do, the no-go list is not specific enough yet.

The five privacy rules

All three of these skills describe a named real person, assembled from their private material. That makes them different from every other skill you will install.

- 1 **Never publish it.** Not to a public repository, not to a skills directory, not to a marketplace. This holds even with names removed, because a profile of how someone works is identifying by construction.
- 2 **Never share it organisation-wide.** On Team and Enterprise plans you can share a skill with a group or the whole company, and the button is right there. This is not one of those. Personal only, or shared with the one named colleague who covers for you.
- 3 **Strip it again at packaging time.** You cleaned the samples once when you built the document. Do it again now. Names, figures, clients, anything about a third person.
- 4 **Tell your principal it exists in this form.** They agreed to a document. Say that it is now something that runs automatically whenever you work on their behalf, and that you will delete it whenever they ask.
- 5 **Delete it when you leave.** It is not yours to take, and it is not the next assistant's to inherit without your principal agreeing again.

THE TEST THAT CATCHES MOST MISTAKES

Before a line goes into a skill you are about to package, ask: *if this file were forwarded to the wrong person tomorrow, which lines would I regret?*

Move those to the annex. If you would regret them being in the annex too, do not write them anywhere.

If your organisation has rules about where material relating to a named individual may be stored, those rules apply to these files. They are closer to a personnel record than to a working document, and they should be handled that way.

ONE MORE, FOR ENTERPRISE

Some organisations scan uploaded skills for malicious content before allowing them to run. Scans usually finish in a minute or two. If a skill is blocked you cannot use it until whatever was flagged is fixed and it is uploaded again. This is a scan of the file, not a review of its contents by a person, but it is worth knowing your file passes through a system before it runs.

When it does not work

WHAT YOU SEE

WHAT IS ACTUALLY WRONG

The skill never fires

Almost always the description, not the body. Rewrite it using the words you actually type. Check for an unquoted colon, which breaks it silently.

It uploaded but nothing happens

It is uploaded but not switched on. Two separate steps. Check the list.

Upload is rejected

The zip contains loose files rather than the folder. Re-zip the folder itself.

The name is rejected

Capitals, spaces, underscores, or the words *claude* or *anthropic*, which are reserved. Lowercase letters, numbers and hyphens only.

Output is a parody of them

Every tendency is being applied at full force. Add the strength labels, and cut the light preferences out of `SKILL.md` entirely.

It invents preferences you never recorded

The instruction to ask rather than guess is missing from the body. Add it explicitly.

Right facts, wrong feel

The samples are not being read. Add the line telling it to read `references/samples.md` when a draft feels wrong.

It was fine, now it is not

Something changed about your principal and the skill did not. New role, new board, new team. Re-run the source skill.

No option to add a skill at all

Your organisation has user-created skills turned off. An administrator setting. Ask.

DO NOT REBUILD WHEN YOU CAN EDIT

Nine times out of ten a wrong result is fixed by one added sentence, not by starting again. Open the skill, highlight the part that is wrong, click "Edit with Claude", and describe the change in plain words.

The whole thing on one page

Every time

- 1 Make a folder named `playbook-jchen`, `voice-jchen` or `linkedin-jchen`.
- 2 Put `SKILL.md` in it. Settings block at the top: `name` in lowercase and hyphens matching the folder, `description` under 200 characters with no colon in it.
- 3 Body holds the rules you want applied every time, in imperative mood, with the numbers and the strength labels kept.
- 4 Long material and verbatim samples go in `references/`, with a line in the body telling Claude when to read them.
- 5 End with the two rules: never change a fact to match a preference, and do not force every pattern in.
- 6 Zip the folder. Customize, Skills, Add, Upload a skill. Switch it on.
- 7 Test against something you have already done, then against something awkward.
- 8 Log every correction. Promote the repeats each quarter.

What never goes in

SKILL	EXCLUDED
<code>playbook-*</code>	Everything in the private annex. Identity, addresses, family, medical, loyalty and payment numbers.
<code>voice-*</code>	Unstripped samples. Names, figures, clients, anything about a third person.
<code>linkedin-*</code>	Their login. Anything not yet public.
All three	Passwords, PINs, recovery codes. And none of these files is ever published, shared organisation-wide, or handed on when you leave.

Where the platform details come from. Character limits, plan requirements, installation steps and sharing behaviour come from Anthropic's own help pages, checked on 31 August 2026. Anthropic changes these. If what you see on screen does not match what is written here, your screen is right and this guide is out of date.

support.claude.com/en/articles/12512198-how-to-create-custom-skills · support.claude.com/en/articles/12512180-use-skills-in-claude

The LinkedIn quotations come from LinkedIn's own help pages and User Agreement, checked on the same date.

The rest is ours, and is not official guidance from anyone. The tiering, the source tags, the naming convention and the maintenance loop are our practice. Adapt them to whatever your organisation already requires for material about a named individual.

[sortedEA.club](#) · [Field guide](#) · [September 2026](#)